

02 – In-Memory Databases

1. Disk Oriented DBMSs

- A. Primary storage : non-volatile storage, organized as fixed length pages
- B. Use buffer pool to cache pages in memory
- C. Concurrency Control :
All txn could stall if it tries to access data that is not in memory.
- D. Logging and recovery :
most use STEAL + NO-FORCE policy. So it must maintain log records to keep track of modification of DB
- E. Overhead
Buffer pool operation, locking and latching is very big overhead
and real work overhead is only 7% of all overhead

2. In-Memory DBMSs

- A. Primary storage : permanently in memory
possible because DRAM prices are low and capacities are high
no need to be organized as fixed length pages
make variable length data blocks and map it with fixed data blocks.
- B. Memory specific indexes' throughput is worse than B+Tree because of cache.
- C. Since all data are in memory,
sequential scan is not better than random access
- D. Logging and recovery
 - i. Still use WAL
 - ii. No need to track LSN, since there are not dirty pages
- E. Overhead
 - i. Overhead for disk I/O is eliminated
 - ii. Network, data copy and move, cache miss are still overhead of system.

3. Concurrency Control Bottlenecks

A. Findings from paper

- i. Allocating unique timestamp can be a bottleneck
- ii. Copying workspace in OCC is bottleneck
- iii. Deadlock detection can be a bottleneck if contention occur a lot
- iv. Because of large abort rate, no-wait and wait die don't work well

B. Bottlenecks

- i. Lock thrashing
 - 1. One txn waits longer to acquire locks, other txn waits longer
- ii. Timestamp allocation
- iii. Memory allocation
 - 1. Copying data on all operations are big overhead