

## 04 – Multi-Version Concurrency Control

### 1. Microsoft Hekaton

#### A. Timestamp management

- i. Each txn is assigned  
BeginTS (when it begins), CommitTS (when it commits)
- ii. Each record has two timestamp field  
BEGIN-TS : CommitTS of txn that created record  
END\_TS : BeginTS of txn that created next version  
or CommitTS of txn that created record

#### B. Version management

- i. Timestamp of uncommitted version's first bit is set to 1 to check it is uncommitted. First bit is excluded when compare it to check whether txn can see the version
- ii. Each txn can read uncommitted version, but cannot update uncommitted version.

#### C. Transaction management

- i. Each txn is at specific "state" :  
ACTIVE, VALIDATING, COMMITTED, TERMINATED
- ii. Txn meta-data  
Read set, Write set, Scan set, Commit dependencies

#### D. Observations

- i. Validations are expensive for Analytical txn
- ii. O2N version chain is not good for OLAP, because of pointer chasing
- iii. Record-level conflict checks can occur false-positive aborts.

## 2. TUM HyPer

### A. Version management

- i. Main data table stores most recent version of records
- ii. Old versions are stored in delta storage, pointed by version vector
- iii. Delta storage is managed per txn, version chain is formed by just pointing next old version in storage.
- iv. Version synopsis is used to skip pointer chasing when there are no other versions but only single main version.

### B. Validation

- i. First writer wins
- ii. Check the redo buffers of txns committed after validating txn began.
- iii. Precision locking is used to validate read set and scan set.

## 3. SAP HANA

### A. Version Management

- i. N2O storage, but store oldest version in main storage.
- ii. Every version has flag to indicate that there are newer versions

## 4. CMU Cicada

### A. Best-Effort Inlining : store meta-data in fixed location

### B. Validation

- i. Contention aware validation
- ii. Early consistency check
- iii. Incremental version search