

05 – Multi-Version Concurrency Control

1. MVCC Deletes

A. Deleted flag

- i. Maintain flag to indicate logical tuple has been deleted
- ii. Separate column / tuple header

B. Tombstone tuple

- i. Create an empty version to indicate deletion
- ii. Use separate pool for it

2. Garbage Collection

A. Reclaimable versions

- i. No active txn can see
- ii. Created by an aborted txn

B. Traditional garbage collection methods cannot handle HTAP well if there are some long running queries

C. Old versions issue

- i. Memory usage increase
- ii. Memory allocator overhead
- iii. Slow read because of longer version chains
- iv. Slow garbage collection because of lots of garbage versions
- v. Poor time-based version locality

D. Design Decisions

- i. Index clean up
Index should clean up key that is not visible anymore
- ii. Version tracking level
 - 1. Tuple level – background vacuuming / cooperative cleaning
 - 2. Txn level – txns keep track of their old versions
 - 3. Epochs – multiple txns are managed in epochs
- iii. Frequency
tradeoffs : more frequent GC
-> efficient memory management, slow down txns.
 - 1. Periodically
run GC at some intervals (fixed vs. adjust on load)
 - 2. Continuously
run GC with txn execution
- iv. Granularity
tradeoffs : more fine granularity
-> more computational overhead for find garbage
 - 1. Single version
high overhead, fine-grained control
 - 2. Group version
less overhead, delay reclamation
 - 3. Tables (If no txns access table, then run GC)
special case
- v. Comparison unit
 - 1. Timestamp – use global minimum timestamp
 - 2. Interval – use timestamp range that are not visible

3. Block Compaction (need to re-study)

A. Deleted tuples

- i. Reuse slots – destroy temporal locality
- ii. Leave slot unoccupied – need some method to fill holes

B. Block compaction

- i. Time since last update
- ii. Time since last access
- iii. Application level semantics