

18 – Sorting Algorithms

1. Background

A. Process

- i. Sort phase
sort the tuples of input relation
- ii. Merge
scan the sorted relation, compare tuples and match the tuples.

B. Hardware-awareness

utilizing CPU, NUMA-aware algorithm implementation,
using SIMD instruction

2. Parallel Sort-Merge Join

A. Process

- i. Partition
partition one input relation and assign them to workers
- ii. Sort
sort the tuples of input relation/partition
- iii. Merge
scan the sorted relation/partition,
compare tuples and match the tuples.

B. Partitioning

- i. Implicit partitioning
data was partitioned on join key when it was loaded
does not always work on OLAP environment
in OLTP environment, most of joins are run on foreign key, so this can work.
- ii. Explicit partitioning
divide outer relation and redistribute it among the different workers.
similar/same with radix partitioning

C. Sort Phase

- i. Create sorted chunks of tuples for both input relations.
- ii. Can use famous sort algorithm like quicksort
need to aware of NUMA and parallel architectures.
- iii. Cache conscious sorting
first, in-register
second, in cache(CPU cache : L1, L2, L3)
last, out of order(if runs of second level exceed the size of caches)
- iv. Sorting networks
fixed wiring paths for lists with the same number of elements
since it does not have branch, and it has limited data dependencies, it is good on modern CPUs.

can be used in level 1 sorting, with SIMD min/max instruction and transposing registers
- v. Bitonic merge network
similar with sorting networks, but merging two sorted registers into single sorted run.
- vi. Multi-way merging
using binotic merge networks, with splitting the process into tasks
similar with merge sort with merging network.
no synchronization, no coordination with other threads.

real DBs do not implement this algorithm.

D. Merge phase

- i. Iterate through outer table and inner table in lockstep and compare join keys.
- ii. Can be done in parallel at the different cores without synchronization.

E. Variant : m-way, m-pass, mpsm

F. Multi-way sort merge join

- i. Sort input relations, partition it by join key, and redistribute partitions among NUMA region(only remote read/write)
- ii. Do local merge join on each NUMA region
since it is partitioned, it does not need remote read/write
hence this will be super fast.

G. Multi-pass sort merge join

- i. Sort input relations on locally on chunks.
- ii. Do global merge join on every pair of NUMA region

H. Massively Parallel Sort Merge join

- i. Range partition outer table, sort each partition locally by different cores
- ii. no partition on inner table, but sort each chunks locally
- iii. do merge sort globally on outer table, locally on inner table.

3. Evaluation

A. Comparison of sort-merge joins

- i. Partition phase's performance is all same on m-way, m-pass, mpsm
- ii. Sorting phase's performance of m-way and m-pass is almost same, but mpsm's sorting phase is a lot slower.
- iii. M-way's merge phase is a lot faster than m-pass's
also, mpsm's merge phase is slowest among the sort-merge joins

B. Scalability of sort-merge joins

- i. M-way is good for scaling on # of threads, but, at hyper-threading area, m-way's scalability degrades.

C. Sort-merge vs. hash

- i. Hash is better than sort-merge at all situation
- ii. Its performance gap narrows when the input relations are huge.