

## 19 – Query Optimizier Overview

### 1. Optimizing

- A. Finding correct query that has lowest model
- B. Proven to be NP-Complete  
So, Optimizer does not always give optimal solution

### 2. Background

#### A. Logical Plan VS. Physical Plan

Optimizer generates a mapping of a logical algebra expression to the optimal equivalent physical algebra expression

ex) access table -> access table with B+Tree index…….,  
choice of join operators…

#### B. Search ARGument ABLE

DBMS can take advantage of an index to speed up the execution of the query.

usually, OLTP queries are sargable. Because they doesn't have table access that much.

#### C. Cost estimation

generate an estimate of the cost of executing a plan for the current state fo the DB

- i. Interactions with other work in DBMS
- ii. Size of result set, data properties
- iii. Choice of algorithms, access methods
- iv. Resource utilization

### 3. Implementation Design Decisions

#### A. Optimization granularity

- i. Single query  
small search space, no reuse result across queries  
just consider what is currently running
- ii. Multiple queries  
large search space, useful for data/intermediate result sharing  
efficient if there are many similar queries

#### B. Optimization timing

- i. Static optimization  
optimize before execution
- ii. Dynamic optimization  
optimize subsequently with execution
- iii. Adaptive optimization  
optimize before execution,  
if performance is lower than threshold, re-optimize query plan and re-execute query.

#### C. Prepared statement (re executing same/similar query)

- i. Reuse last plan
- ii. Re-optimize
- iii. Multiple plans
- iv. Average plan

#### D. Plan stability

- i. Hints  
DBA can provide hints to the optimizer.(some constraints on optimizing)
- ii. Fixed optimizer versions  
set the optimizer version number and migrate queries one-by-one to the new optimizer
- iii. Backwards compatible plans  
save query plan from old version and provide it to the new DBMS

#### E. Search termination

- i. Wall-clock time  
stop optimizer at some length of time
- ii. Cost threshold  
stop optimizer when it finds a plan that has lower cost than some threshold
- iii. Exhaustion  
stop when there are no more enumerations of the target plan.

### 4. Optimizer Search Strategies

#### A. Heuristics

- i. rules
  1. Perform most restrictive selection early
  2. Perform all selections before joins
  3. Join ordering based on cardinality
- ii. Pros & cons
  1. Easy to implement and debug, works reasonably well.
  2. If operators have complex inter-dependencies, it'll not work.

## B. Heuristics + cost-based join order search

### i. Rules

1. Use static rules to perform initial optimization
2. Use dynamic programming to determine the best join order

### ii. Pros & cons

1. Can find a reasonable plan without an exhaustive search
2. Same problem with heuristic only approach
3. Left-deep join trees are not always optimal

## C. Randomized algorithms

### i. Rules

1. Select random query plan until cost threshold is reached
2. Simulated annealing : start with initial query plan which is generated by heuristic only approach, do random change of operator and accept when cost is reduced.

### ii. Pros & Cons

1. can avoid local minimum trap
2. low memory overhead
3. difficult to determine why the DBMS may have chosen a plan

#### D. Stratified search

##### i. Rules

1. Rewrite the logical query plan using transformation rules
2. Then perform a cost-based search to map the logical plan to a physical plan

##### ii. Pros & cons

1. Works well in practice
2. Difficult to assign priorities to transformations
3. Rules maintenance is big pain

#### E. Unified search

##### i. Rules

1. Unify the notion of logical-logical and logical-physical transformation
2. Use memoization to reduce redundant work since it generates many transformations