

## 21 – Query Optimizer Implementation

### 1. Bad query plan

- A. Common problem :  
incorrect join orderings because of inaccurate cardinality estimation
- B. Reason of plans getting worse than expectation :  
estimations are based on static summarizations  
can be solved by adaptive query optimization

### 2. Adaptive query optimization

- A. Modify the execution behavior of a query by generating multiple plans for it
- B. Use information collected during query execution.  
(affected by previous query runs)
- C. Approach
  - i. Modify future invocation
  - ii. Replan current invocations
  - iii. Plan pivot points

### 3. Reversion-Based Plan Correction

- A. Check new plan to previously executed plan  
if new plan is not better than previous one, change query plan with previous one.
- B. Can recover from wrong estimated cost because comparison of query plans are done by actual cost.  
even if first estimation goes wrong, future invocation won't suffer by this estimation

#### 4. Plan Stitching

- A. Using sub-plan of previously executed plan to build plan to compare with new plan
- B. Use heuristics to identify equivalent sub-plans
- C. Use OR-operators to indicate alternative paths through the plan  
AND-OR tree will be DAG structure
- D. Perform bottom-up search to select the cheapest sub-plan for each OR nodes.
- E. Redshift  
DBMS caches subplans(compiled one), combines these at runtime for new queries. (to avoid the compilation cost)

#### 5. Quickstep

#### 6. Parametric optimization

#### 7. Proactive re-optimization