

07 – Tree Indexes 1

1. Table index

- A. Synchronized replica of a subset of a table's attributes that are organized for access using a subset of attributes
- B. Storage & Maintenance Overhead

2. B+Tree

- A. Self balancing tree structure
- B. Keeps data sorted
- C. $O(\log N)$ complexity
- D. B+Tree has records only in leaf nodes
- E. Advantage
 - i. Clustered Index
 - 1. Records are sorted in order of primary key (hash table is not)
 - ii. Selection Condition
 - 1. If query has data for some attributes, Tree structure can search that data (hash table need all of attributes)

3. Design Decisions

- A. Node Size : Disk slow? Enlarge node size!
- B. Delayed Merge : Merge may occur more re-organization
- C. Variable length keys
 - i. Pointers
 - ii. Various size of nodes
 - iii. Padding
 - iv. Key mapping
- D. Non-unique keys : Duplicate or List?
- E. Intra-Node Search : Linear or Binary (or Interpolate)

4. Optimization

- A. Prefix Compression - common prefix only stored once
- B. Suffix Truncation - keys are only used as traffic, So, only needed part stored
- C. Bulk Insert - sort data first, construct tree by bottom up logic
- D. Pointer Swizzling -
if page is pinned, don't lookup the page table, just get pointer

5. What is it?

- A. Pointer Swizzling (I didn't get it well)

6. Introduced Papers

- A. Efficient Locking for Concurrent Operations on B-Trees
- B. Modern B-Tree Techniques